

# Class 8: Breast Cancer Mini Project

Bryn Baxter (A69038039)

## Table of contents

Tidy to remove the diagnosis column. . . . .	4
Cluster the data set . . . . .	4
Principal Component Analysis (PCA) . . . . .	6
Breast Cancer data PCA . . . . .	9

Today we will practice applying our PCA and clustering methods from last class on some simple FNA breast cancer data.

Lets get the data into R...

```
wisc.df <- read.csv("WisconsinCancer.csv", row.names=1)
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	
842302	M	17.99	10.38	122.80	1001.0	
842517	M	20.57	17.77	132.90	1326.0	
84300903	M	19.69	21.25	130.00	1203.0	
84348301	M	11.42	20.38	77.58	386.1	
84358402	M	20.29	14.34	135.10	1297.0	
843786	M	12.45	15.70	82.57	477.1	
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean		
842302	0.11840	0.27760	0.3001	0.14710		
842517	0.08474	0.07864	0.0869	0.07017		
84300903	0.10960	0.15990	0.1974	0.12790		
84348301	0.14250	0.28390	0.2414	0.10520		
84358402	0.10030	0.13280	0.1980	0.10430		
843786	0.12780	0.17000	0.1578	0.08089		
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se	
842302	0.2419	0.07871	1.0950	0.9053	8.589	
842517	0.1812	0.05667	0.5435	0.7339	3.398	

84300903	0.2069		0.05999	0.7456	0.7869	4.585
84348301	0.2597		0.09744	0.4956	1.1560	3.445
84358402	0.1809		0.05883	0.7572	0.7813	5.438
843786	0.2087		0.07613	0.3345	0.8902	2.217
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302	153.40	0.006399	0.04904	0.05373		0.01587
842517	74.08	0.005225	0.01308	0.01860		0.01340
84300903	94.03	0.006150	0.04006	0.03832		0.02058
84348301	27.23	0.009110	0.07458	0.05661		0.01867
84358402	94.44	0.011490	0.02461	0.05688		0.01885
843786	27.19	0.007510	0.03345	0.03672		0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst		
842302	0.03003		0.006193	25.38		17.33
842517	0.01389		0.003532	24.99		23.41
84300903	0.02250		0.004571	23.57		25.53
84348301	0.05963		0.009208	14.91		26.50
84358402	0.01756		0.005115	22.54		16.67
843786	0.02165		0.005082	15.47		23.75
	perimeter_worst	area_worst	smoothness_worst	compactness_worst		
842302	184.60	2019.0		0.1622		0.6656
842517	158.80	1956.0		0.1238		0.1866
84300903	152.50	1709.0		0.1444		0.4245
84348301	98.87	567.7		0.2098		0.8663
84358402	152.20	1575.0		0.1374		0.2050
843786	103.40	741.6		0.1791		0.5249
	concavity_worst	concave.points_worst	symmetry_worst			
842302	0.7119		0.2654			0.4601
842517	0.2416		0.1860			0.2750
84300903	0.4504		0.2430			0.3613
84348301	0.6869		0.2575			0.6638
84358402	0.4000		0.1625			0.2364
843786	0.5355		0.1741			0.3985
	fractal_dimension_worst					
842302		0.11890				
842517		0.08902				
84300903		0.08758				
84348301		0.17300				
84358402		0.07678				
843786		0.12440				

Q1. How many samples/patients are in this dataset?

There are 569 samples in this data set.

Q2. How many cancer/non-cancer diagnosis samples are there?

```
sum(wisc.df$diagnosis == "M")
```

```
[1] 212
```

The `table()` function is a super useful utility for counting the number of observations.

```
table(wisc.df$diagnosis)
```

```
  B   M  
357 212
```

Q3. How many dimensions are there in this dataset?

```
ncol(wisc.df)
```

```
[1] 31
```

Q4. How many columns are suffixed with “\_mean”?

```
colnames(wisc.df)
```

```
[1] "diagnosis"           "radius_mean"  
[3] "texture_mean"        "perimeter_mean"  
[5] "area_mean"           "smoothness_mean"  
[7] "compactness_mean"    "concavity_mean"  
[9] "concave.points_mean" "symmetry_mean"  
[11] "fractal_dimension_mean" "radius_se"  
[13] "texture_se"          "perimeter_se"  
[15] "area_se"             "smoothness_se"  
[17] "compactness_se"      "concavity_se"  
[19] "concave.points_se"   "symmetry_se"  
[21] "fractal_dimension_se" "radius_worst"  
[23] "texture_worst"       "perimeter_worst"  
[25] "area_worst"          "smoothness_worst"  
[27] "compactness_worst"   "concavity_worst"  
[29] "concave.points_worst" "symmetry_worst"  
[31] "fractal_dimension_worst"
```

The `grep()` function can help us find pattern matches here:

```
x <- grep("_mean", colnames(wisc.df))  
  
length(x)
```

```
[1] 10
```

### **Tidy to remove the diagnosis column.**

Lets get rid of the diagnosis column.

Save a vector of this expert diagnosis for later and remove it from the data to undergo clustering, PCA. etc...

```
diagnosis <- wisc.df$diagnosis
```

```
wisc.data <- wisc.df[,-1]  
dim(wisc.data)
```

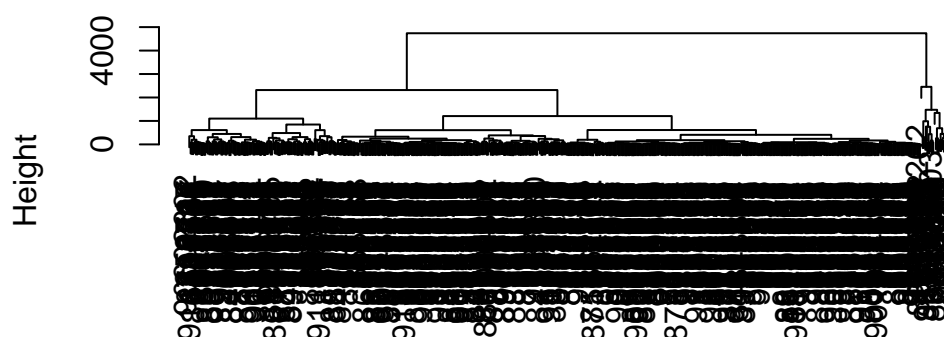
```
[1] 569 30
```

### **Cluster the data set**

Lets try a `hclust()`.

```
hc.row <- hclust(dist(wisc.data))  
plot(hc.row)
```

## Cluster Dendrogram



```
dist(wisc.data)
hclust(*, "complete")
```

To get some clusters out of this, we can cut the tree at a given height.

```
grps <- cutree(hc.row, h=4000)
table(grps)
```

```
grps
 1    2
549  20
```

To see the correspondence of our cluster `grps` with the expert `dianosis` I can use `table()`:

```
table(grps, diagnosis)
```

```
      diagnosis
grps   B    M
 1  357 192
 2    0   20
```

That is not that a useful clustering result....

## Principal Component Analysis (PCA)

Scaling data before analysis is often critical.

Side Note: the default for `prcomp()` is `scale=FALSE`.

There is a dataset in R called `mtcars` which has loads of numbers about old cars.

```
head(mtcars)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

```
colMeans(mtcars)
```

mpg	cyl	disp	hp	drat	wt	qsec
20.090625	6.187500	230.721875	146.687500	3.596563	3.217250	17.848750
vs	am	gear	carb			
0.437500	0.406250	3.687500	2.812500			


```
apply(mtcars, 2, sd)
```

mpg	cyl	disp	hp	drat	wt
6.0269481	1.7859216	123.9386938	68.5628685	0.5346787	0.9784574
qsec	vs	am	gear	carb	
1.7869432	0.5040161	0.4989909	0.7378041	1.6152000	


```
pc.noscale <- prcomp(mtcars, scale=F)
pc.scale <- prcomp(mtcars, scale=T)
```

Lets look at the loadings first:

```
library(ggplot2)
ggplot(pc.noscale$rotation) +
  aes(PC1, rownames(pc.noscale$rotation))+
  geom_col()
```




```
ggplot(pc.scale$rotation) +
  aes(PC1, rownames(pc.scale$rotation))+
  geom_col()
```




The main PCA result figure is often called a “score plot” or “PC plot” or “PC1 vs. PC2 plot”

```
ggplot(pc.noscale$x) +  
  aes(PC1, PC2, label=rownames(pc.noscale$x)) +  
  geom_point()+  
  geom_label()
```



```
ggplot(pc.scale$x) +  
  aes(PC1, PC2, label=rownames(pc.scale$x)) +  
  geom_point()+  
  geom_label()
```





```
x <- scale(mtcars)
round(colMeans(x))
```

```
mpg   cyl  disp    hp  drat    wt   qsec    vs    am  gear  carb
0     0     0     0     0     0     0     0     0     0     0
```

```
round(apply(x, 2, sd))
```

```
mpg   cyl  disp    hp  drat    wt   qsec    vs    am  gear  carb
1     1     1     1     1     1     1     1     1     1     1
```

**Key Point:** Generally we want to “scale” our data before analysis to avoid being mis-lead due to your data having different measurement units.

## Breast Cancer data PCA

We will scale our data.

```
pca <- prcomp(wisc.data, scale=T)
```

See how well we are doing :

```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010

	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335

	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966


	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997

	PC29	PC30
Standard deviation	0.02736	0.01153
Proportion of Variance	0.00002	0.00000
Cumulative Proportion	1.00000	1.00000

Our PC plot:

```
ggplot(pca$x) +  
  aes(PC1, PC2, col=diagnosis) +  
  geom_point()+  
  xlab("PC1 (44.3%)" )+  
  ylab("PC2 (18.9%)")
```



Each point represents a sample and its measured cell characteristics, the general idea is that samples with similar characteristics will cluster together.

PCA is a method that takes a data set with a lot of dimensions and flattens it down to 2 or 3 dimensions so we can look at it.

Q5. How many PCs capture 80% of the original variance?

```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966


	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997

	PC29	PC30
Standard deviation	0.02736	0.01153
Proportion of Variance	0.00002	0.00000
Cumulative Proportion	1.00000	1.00000

5 Pcs.

```
plot(pca)
```



Q6. Use ggplot to plot a “scree-plot” of the variance captured per PC.

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation" "center"    "scale"     "x"

$class
[1] "prcomp"
```

We can extract the sdev and figure out the variance.

```
v <- pca$sdev^2  
round(v)
```

```
[1] 13  6  3  2  2  1  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
[26]  0  0  0  0  0
```

```
#to get total variance  
sum(v)
```

```
[1] 30
```

The proportion of variance captured in PC

```
round(v/sum(v), 2)
```

```
[1] 0.44 0.19 0.09 0.07 0.05 0.04 0.02 0.02 0.01 0.01 0.01 0.01 0.01 0.01 0.00  
[16] 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00
```

Cumulative variance captured

```
cumsum(v/sum(v))
```

```
[1] 0.4427203 0.6324321 0.7263637 0.7923851 0.8473427 0.8875880 0.9100953  
[8] 0.9259825 0.9398790 0.9515688 0.9613660 0.9700714 0.9781166 0.9833503  
[15] 0.9864881 0.9891502 0.9911302 0.9928841 0.9945334 0.9955720 0.9965711  
[22] 0.9974858 0.9982971 0.9988990 0.9994150 0.9996876 0.9999176 0.9999706  
[29] 0.9999956 1.0000000
```

```
which(cumsum(v/sum(v))>0.8)
```


```
[1]  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29  
[26] 30
```

```
#to see the first one, use [1]
```

```
library(factoextra)
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

```
fviz_eig(pca, addlabels = T)
```




## Combine PCA and clustering

We saw earlier that clustering the raw data alone did not look that useful.

We can use our new PC variables (our PCs) as a basis for clustering. Use our `$x` PC scores and cluster in the PC1-2 subspace

```
hc.pca <- hclust(dist(pca$x[,1:2]), method="ward.D2")  
plot(hc.pca)  
abline(h=60, col="blue")
```


## Cluster Dendrogram



```
dist(pca$x[, 1:2])  
hclust (*, "ward.D2")
```

Q7. Does your clustering help separate cancer from non-cancer samples (i.e. diagnosis “M vs. B”)

```
grps.1 <- cutree(hc.pca, h=60)  
plot(grps.1)
```



```
table(grps.1, diagnosis)
```

```
      diagnosis
grps.1  B    M
1      18 177
2     339  35
```

Positive cancer samples “M” Negative cancer samples “B”

True our cluster/grp1 False out cluster/grp2

Q8. How many True positives (TP) do we have?

Q9. How many false positives (FP) do we have?

Sensitivity  $TP/(TP+FN)$ . Specificity  $TN/(TN+FN)$ .

##Prediction with our PCA model

We can think new data (in this case from UofM) and project it onto our new variables PCs.

Read UofM data



```
url <- "https://tinyurl.com/new-samples-CSV"  
new <- read.csv(url)
```

Projection

```
npc <- predict(pca, newdata=new)
```

Base R plot

```
plot(pca$x[,1:2], col=grps.1)  
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)  
text(npc[,1], npc[,2], c(1,2), col="white")
```

